

```
""" DL ASSIGNMENT
```

```
ML | Credit Card Fraud Detection
```

```
-> The challenge is to recognize fraudulent credit card transactions  
so that the customers of credit card companies  
are not charged for items that they did not purchase.
```

```
Main challenges involved in credit card fraud detection are:
```

- 1.Enormous Data is processed every day and the model build must be fast enough to respond to the scam in time.
- 2.Imbalanced Data i.e most of the transactions (99.8%) are not fraudulent which makes it really hard for detecting the fraudulent ones
- 3.Data availability as the data is mostly private.
- 4.Misclassified Data can be another major issue, as not every fraudulent transaction is caught and reported.
- 5.Adaptive techniques used against the model by the scammers.

```
"""
```

```
from google.colab import files
```

```
uploaded = files.upload()
```



Choose Files creditcard.csv

- **creditcard.csv**(text/csv) - 150828752 bytes, last modified: 9/25/2022 - 100% done
Saving creditcard.csv to creditcard.csv

```
#Setup
```

```
#We will be using TensorFlow 1.2 and Keras 2.0.4.
```

```
import pandas as pd
```

```
import numpy as np
```

```
import pickle
```

```
import matplotlib.pyplot as plt
```

```
from scipy import stats
```

```
import tensorflow as tf
```

```
import seaborn as sns
```

```
from pylab import rcParams
```

```
from sklearn.model_selection import train_test_split
```

```
from keras.models import Model, load_model
```

```
from keras.layers import Input, Dense
```

```
from keras.callbacks import ModelCheckpoint, TensorBoard
```

```
from keras import regularizers
```

```
%matplotlib inline
```

```
sns.set(style='whitegrid', palette='muted', font_scale=1.5)
```

```
rcParams['figure.figsize'] = 14, 8
```

```
RANDOM_SEED = 42
```

```
LABELS = ["Normal", "Fraud"]
```

```
"""we will train an Autoencoder Neural Network (implemented in Keras) in unsupervised (or
The trained model will be evaluated on pre-labeled and anonymized dataset."""
```

```
#Loading the data
```

```
#The dataset contains data about credit card transactions that occurred during a period of
```

```
"""All variables in the dataset are numerical. The data has been transformed using PCA tra
changed are Time and Amount. Time contains the seconds elapsed between each transaction an
```

```
import pandas as pd
```

```
import io
```

```
df = pd.read_csv(io.StringIO(uploaded['creditcard.csv'].decode('utf-8')))
```

```
print(df)
```

	Time	V1	V2	V3	V4	V5	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
...	
284802	172786.0	-11.881118	10.071785				

```
0      0
1      0
2      0
3      0
4      0
...    ...
284802  0
284803  0
284804  0
284805  0
284806  0
```

```
[284807 rows x 31 columns]
```

```
#Exploration
df.shape
```


```
(284807, 31)
```

```
#31 columns, 2 of which are Time and Amount. The rest are output from the PCA transformati
```

```
df.isnull().values.any()
```

```
False
```

```
count_classes = pd.value_counts(df['Class'], sort = True)
count_classes.plot(kind = 'bar', rot=0)
plt.title("Transaction class distribution")
plt.xticks(range(2), LABELS)
plt.xlabel("Class")
plt.ylabel("Frequency");
```



#We have a highly imbalanced dataset on our hands. Normal transactions overwhelm the fraud

```
≈
```

```
frauds = df[df.Class == 1]
normal = df[df.Class == 0]
```

```
100000
```

```
frauds.shape
```

```
(492, 31)
```

```
normal.shape
```

```
(284315, 31)
```

#How different are the amount of money used in different transaction classes?

```
frauds.Amount.describe()
```

```
count    492.000000
mean     122.211321
std      256.683288
min       0.000000
25%       1.000000
50%       9.250000
75%      105.890000
max      2125.870000
Name: Amount, dtype: float64
```

```
normal.Amount.describe()
```

```
count    284315.000000
mean       88.291022
std       250.105092
min        0.000000
25%        5.650000
50%       22.000000
75%       77.050000
max     25691.160000
Name: Amount, dtype: float64
```

#Let's have a more graphical representation:


```
f, (ax1, ax2) = plt.subplots(2, 1, sharex=True)
f.suptitle('Amount per transaction by class')
```

```
bins = 50
```

```
ax1.hist(frauds.Amount, bins = bins)
ax1.set_title('Fraud')
```

```
ax2.hist(normal.Amount, bins = bins)
ax2.set_title('Normal')
```

```
plt.xlabel('Amount ($)')
plt.ylabel('Number of Transactions')
plt.xlim((0, 20000))
plt.yscale('log')
plt.show();
```



#Do fraudulent transactions occur more often during certain time?


```
f, (ax1, ax2) = plt.subplots(2, 1, sharex=True)
f.suptitle('Time of transaction vs Amount by class')
```

```
ax1.scatter(frauds.Time, frauds.Amount)
ax1.set_title('Fraud')
```

```
ax2.scatter(normal.Time, normal.Amount)
```

```
ax2.set_title('Normal')
```

```
plt.xlabel('Time (in Seconds)')
plt.ylabel('Amount')
plt.show()
```



```
"""Autoencoders
```

Autoencoders can seem quite bizarre at first. The job of those models is to predict the in
More specifically, let's take a look at Autoencoder Neural Networks. This autoencoder tries

While trying to do just that might sound trivial at first, it is important to note that we
This can be done by limiting the number of hidden units in the model. Those kind of autoe
"""

```
"""
```

```
-----Reconstruction error-----
```

We optimize the parameters of our Autoencoder model in such way that a special kind of err

```
"""
```

```
-----Preparing the data-----
```

First, let's drop the Time column (not going to use it) and use the scikit's StandardScale

```

from sklearn.preprocessing import StandardScaler

data = df.drop(['Time'], axis=1)

data['Amount'] = StandardScaler().fit_transform(data['Amount'].values.reshape(-1, 1))

X_train, X_test = train_test_split(data, test_size=0.2, random_state=RANDOM_SEED)
X_train = X_train[X_train.Class == 0]
X_train = X_train.drop(['Class'], axis=1)

y_test = X_test['Class']
X_test = X_test.drop(['Class'], axis=1)

X_train = X_train.values
X_test = X_test.values

X_train.shape

(227451, 29)

"""
-----Building the model-----
Our Autoencoder uses 4 fully connected layers with 14, 7, 7 and 29 neurons respectively. T
Additionally, L1 regularization will be used during training:"""

input_dim = X_train.shape[1]
encoding_dim = 14

input_layer = Input(shape=(input_dim, ))

encoder = Dense(encoding_dim, activation="tanh",
                activity_regularizer=regularizers.l1(10e-5))(input_layer)
encoder = Dense(int(encoding_dim / 2), activation="relu")(encoder)

decoder = Dense(int(encoding_dim / 2), activation='tanh')(encoder)
decoder = Dense(input_dim, activation='relu')(decoder)

autoencoder = Model(inputs=input_layer, outputs=decoder)

"""Let's train our model for 100 epochs with a batch size of 32 samples and save the best
The ModelCheckpoint provided by Keras is really handy for such tasks. Additionally, the t

nb_epoch = 100
batch_size = 32

autoencoder.compile(optimizer='adam',
                    loss='mean_squared_error',
                    metrics=['accuracy'])

```

```

checkpointer = ModelCheckpoint(filepath="model.h5",
                               verbose=0,
                               save_best_only=True)
tensorboard = TensorBoard(log_dir='./logs',
                           histogram_freq=0,
                           write_graph=True,
                           write_images=True)

history = autoencoder.fit(X_train, X_train,
                          epochs=nb_epoch,
                          batch_size=batch_size,
                          shuffle=True,
                          validation_data=(X_test, X_test),
                          verbose=1,
                          callbacks=[checkpointer, tensorboard]).history

```

7108/7108 [=====] - 10s 2ms/step - loss: 0.7343 - accurac
 Epoch 73/100
 7108/7108 [=====] - 15s 2ms/step - loss: 0.7344 - accurac
 Epoch 74/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7343 - accurac
 Epoch 75/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7340 - accurac
 Epoch 76/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7342 - accurac
 Epoch 77/100
 7108/7108 [=====] - 15s 2ms/step - loss: 0.7342 - accurac
 Epoch 78/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7338 - accurac
 Epoch 79/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7341 - accurac
 Epoch 80/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7339 - accurac
 Epoch 81/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7338 - accurac
 Epoch 82/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7340 - accurac
 Epoch 83/100
 7108/7108 [=====] - 15s 2ms/step - loss: 0.7336 - accurac
 Epoch 84/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7338 - accurac
 Epoch 85/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7338 - accurac
 Epoch 86/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7335 - accurac
 Epoch 87/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7337 - accurac
 Epoch 88/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7337 - accurac
 Epoch 89/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7334 - accurac
 Epoch 90/100
 7108/7108 [=====] - 15s 2ms/step - loss: 0.7333 - accurac
 Epoch 91/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7334 - accurac
 Epoch 92/100
 7108/7108 [=====] - 16s 2ms/step - loss: 0.7336 - accurac
 Epoch 93/100


```

7108/7108 [=====] - 15s 2ms/step - loss: 0.7335 - accurac
Epoch 94/100
7108/7108 [=====] - 16s 2ms/step - loss: 0.7333 - accurac
Epoch 95/100
7108/7108 [=====] - 16s 2ms/step - loss: 0.7334 - accurac
Epoch 96/100
7108/7108 [=====] - 21s 3ms/step - loss: 0.7334 - accurac
Epoch 97/100
7108/7108 [=====] - 17s 2ms/step - loss: 0.7337 - accurac
Epoch 98/100
7108/7108 [=====] - 17s 2ms/step - loss: 0.7335 - accurac
Epoch 99/100
7108/7108 [=====] - 19s 3ms/step - loss: 0.7334 - accurac
Epoch 100/100
7108/7108 [=====] - 17s 2ms/step - loss: 0.7334 - accurac

```


```
autoencoder = load_model('model.h5')
```

```
#Evaluation
```

```

plt.plot(history['loss'])
plt.plot(history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper right');


```



```
predictions = autoencoder.predict(X_test)
```

```
mse = np.mean(np.power(X_test - predictions, 2), axis=1)
error_df = pd.DataFrame({'reconstruction_error': mse,
                        'true_class': y_test})
```

```
error_df.describe()
```


	reconstruction_error	true_class	
count	56962.000000	56962.000000	
mean	0.763890	0.001720	
std	3.365745	0.041443	
min	0.051223	0.000000	
25%	0.283433	0.000000	
50%	0.436176	0.000000	
75%	0.660284	0.000000	
max	256.665093	1.000000	

```
#Reconstruction error without fraud
```

```
fig = plt.figure()
```

```
ax = fig.add_subplot(111)
```

```
normal_error_df = error_df[(error_df['true_class']== 0) & (error_df['reconstruction_error']  
_ = ax.hist(normal_error_df.reconstruction_error.values, bins=10)
```




#Reconstruction error with fraud

```
fig = plt.figure()
```

```
ax = fig.add_subplot(111)
```

```
fraud_error_df = error_df[error_df['true_class'] == 1]
```

```
_ = ax.hist(fraud_error_df.reconstruction_error.values, bins=10)
```




```
from sklearn.metrics import (confusion_matrix, precision_recall_curve, auc,
                             roc_curve, recall_score, classification_report, f1_score,
                             precision_recall_fscore_support)
```

"""ROC curves are very useful tool for understanding the performance of binary classifiers
We have a very imbalanced dataset. Nonetheless, let's have a look at our ROC curve:"""

```
fpr, tpr, thresholds = roc_curve(error_df.true_class, error_df.reconstruction_error)
roc_auc = auc(fpr, tpr)
```

```
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, label='AUC = %.4f'% roc_auc)
plt.legend(loc='lower right')
plt.plot([0,1],[0,1], 'r--')
plt.xlim([-0.001, 1])
```


```
plt.ylim([0, 1.001])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show();
```



""The ROC curve plots the true positive rate versus the false positive rate, over different thresholds. Basically, we want the blue line to be as close as possible to the upper left corner.""


""Let's return to our example from Information Retrieval. High recall but low precision means many returned results are not relevant. When precision is high but recall is low we have the opposite - few returned results with high relevance. Ideally, you would want high precision and high recall - many results with that are highly relevant.""

```
precision, recall, thresholds = precision_recall_curve(error_df.true_class, error_df.reconstruction)
plt.plot(recall, precision, 'b', label='Precision-Recall curve')
plt.title('Recall vs Precision')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.show()
```




""A high area under the curve represents both high recall and high precision, where high and high recall relates to a low false negative rate. High scores for both show that the c a majority of all positive results (high recall)."""

```
plt.plot(th, precision[1:], 'b', label='Threshold-Precision curve')
plt.title('Precision for different threshold values')
plt.xlabel('Threshold')
plt.ylabel('Precision')
plt.show()
```



#You can see that as the reconstruction error increases our precision rises as well. Let's

```
plt.plot(th, recall[1:], 'b', label='Threshold-Recall curve')
plt.title('Recall for different threshold values')
plt.xlabel('Reconstruction error')
plt.ylabel('Recall')
plt.show()
```



```
"""
```

```
-----Prediction-----
```


Our model is a bit different this time. It doesn't know how to predict new values. But we is normal or fraudulent, we'll calculate the reconstruction error from the transaction da If the error is larger than a predefined threshold, we'll mark it as a fraud (since our m Let's pick that value:""

```
threshold = 2.9
```

#And see how well we're dividing the two types of transactions:

```
groups = error_df.groupby('true_class')
fig, ax = plt.subplots()

for name, group in groups:
    ax.plot(group.index, group.reconstruction_error, marker='o', ms=3.5, linestyle='',
            label= "Fraud" if name == 1 else "Normal")
ax.hlines(threshold, ax.get_xlim()[0], ax.get_xlim()[1], colors="r", zorder=100, label='Th')
ax.legend()
plt.title("Reconstruction error for different classes")
plt.ylabel("Reconstruction error")
plt.xlabel("Data point index")
plt.show();
```

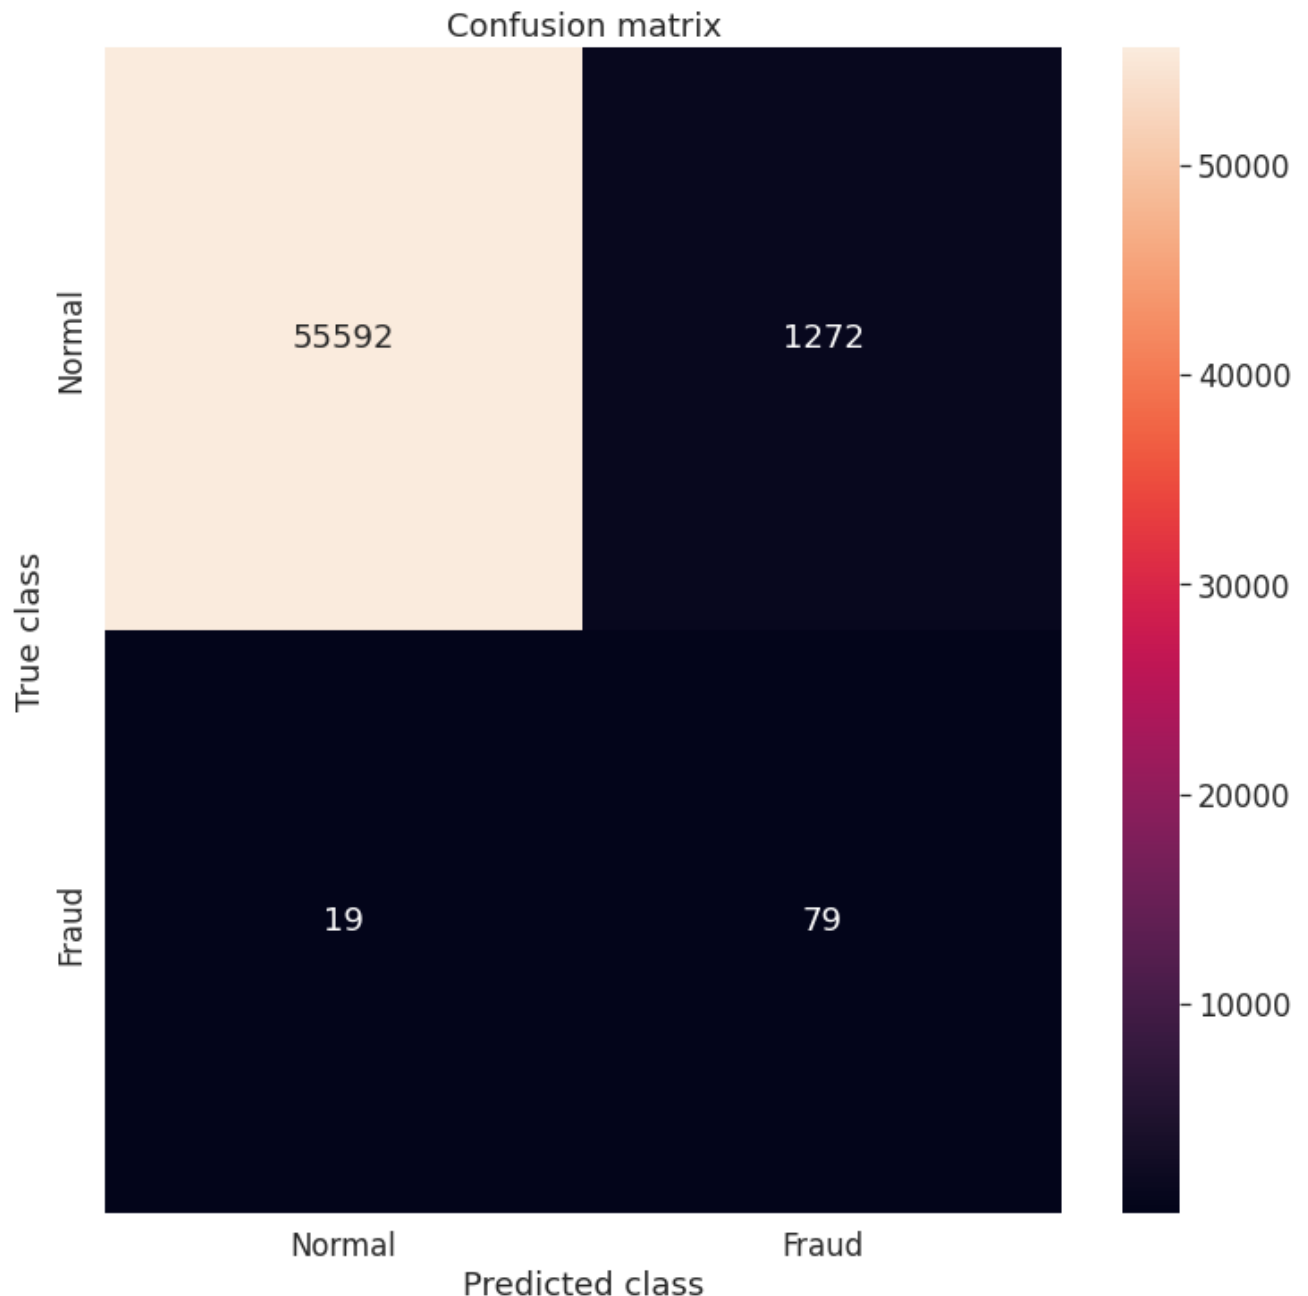


Let's have a look at the confusion matrix:

```
y_pred = [1 if e > threshold else 0 for e in error_df.reconstruction_error.values]
conf_matrix = confusion_matrix(error_df.true_class, y_pred)

plt.figure(figsize=(12, 12))
sns.heatmap(conf_matrix, xticklabels=LABELS, yticklabels=LABELS, annot=True, fmt="d");
plt.title("Confusion matrix")
plt.ylabel('True class')
```

```
plt.xlabel('Predicted class')
plt.show()
```



```
"""
```

```
-----
```

Our model seems to catch a lot of the fraudulent cases. Of course, there is a catch (see w
The number of normal transactions classified as frauds is really high. Is this really a p
depending on the problem. That one is up to you.

```
-----Conclusion-----
```

We've created a very simple Deep Autoencoder in Keras that can reconstruct what non fraudu

Keras gave us very clean and easy to use API to build a non-trivial Deep Autoencoder.
You can search for TensorFlow implementations and see for yourself how much boilerplate yo

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 9:32 PM ● ✕